



Felix Zhang

✉ felixfzhang@cs.toronto.edu |  [ff-zhang](https://github.com/ff-zhang) |  [felixfzhang](https://www.linkedin.com/in/felixfzhang)

EDUCATION

Honours Bachelor of Science, University of Toronto

Sept. 2021 – June 2025

Specialist in Computer Science, Major in Mathematics

3.95/4.0 cGPA


RESEARCH EXPERIENCE

Research Assistant; Far Data Lab, University of Toronto

Sept. 2024 – Present

Supervisor: Prof. Qizhen Zhang

- Investigated the applicability of offloading data processing onto data processing units within data centers
- Parallelized the execution of *Monodepth2* in **Python** and **C++** with the latter achieving linear performance scaling with the number of threads
- Worked on a team of 6 to implement a prefetcher with DPUs that intercepts network traffic to predict the data that will be requested by a tiered key-value store

Research Assistant; University of Toronto 

March 2024 – Dec. 2024

Supervisor: Prof. Jack Sun


- Worked on a team of 11 to implement a pedagogical kernel *KidneyOS* in **Rust** to be used in an introductory operating systems course with **500+** students annually
- Enabled thread creation and destruction, multi-threading, pre-emptive scheduling within the thread system
- Led a team of 3 to implement POSIX-compatible syscalls and add support running user-space executables

Research Assistant; PRISM Lab, Bloorview Research Institute

June 2024 – Aug. 2024

Supervisors: Erica Floreani and Prof. Tom Chau; Funded by: FUSRP


- Curated deep-learning models from the literature on denoising electroencephalogram (EEG) data in a team of 4 and benchmarked them on the *EEGDenoiseNet* dataset
- Investigated the applicability of end-to-end transformer models to denoise EEG signals and the impact of using signals' time-frequency representation as input on model performance

Machine Learning Researcher; BMO Lab, University of Toronto 

July 2023 – May 2024

Supervisor: Prof. David Rokeby; Funded by: Work Study Program

- Applied forward dynamics in real-time on motion-capture data using **MuJoCo**, providing joint-level control of the model and the option to extract physical data using inverse dynamics
- Used imitation learning to enable humanoid models to copy movements from motion capture suits in real-time

Research Assistant; Biological Physics Group, University of Toronto 

May 2023 – May 2024

Supervisor: Prof. Anton Zilman; Funded by: Work Study Program

- Implemented a data pre-processing pipeline which processes raw cytokine data and extracts integral features
- Built a feed-forward network in **PyTorch** that predicts the cytokine dynamics of T cells in response to antigens
- On a team of 4, showed two variables are sufficient to determine cytokine concentrations because our model predicted the correct output concentration with **0.01%** error using a bottleneck layer with 2 neurons

Research Assistant; Physics Education Group, University of Toronto 

May 2022 – Sept. 2022

Supervisor: Prof. Carolyn Sealfon

- Created a dataset of **~11 000** sentences from student feedback which labels whether they contain suggestions
- Compared the effectiveness of statistical and deep-learning classifiers at identifying suggestions using **scikit-learn** and **TensorFlow** respectively
- Demonstrated the efficacy of a BERT classifier at addressing this problem with it achieving an F_1 score of **0.823**

INDUSTRY EXPERIENCE

ML Runtime Engineer; Cerebras Systems

May 2024 – Present

- Implemented a runtime virtual memory system in **C++** with a team of 3 which pre-emptively loads data before it is accessed, allowing off-chip memory to be used for the first time with a **10%** performance penalty
- Added support for network storage in the paging system with remote direct memory access, providing **100 GB/s** read and write speeds with **10 μ s** latency to multiple remote servers
- Enabled logging of all network operations and the ability to later replay the logged operations, decreasing the time to recreate stalls and timeouts by over **80%**

AWARDS & SCHOLARSHIPS

Fields Undergraduate Summer Research Program (\$3 800), Fields Institute	June – Aug. 2024
Louis Savlov Scholarship in Sciences And Humanities (\$1 000), University of Toronto	Nov. 2023 – Jan. 2025
Dean's List Scholar , University of Toronto	Jan. 2022 – Present
Second Malcom Wallace Scholarship (\$5 000), University of Toronto	Sept. 2021 – Oct. 2024
University of Toronto Scholar (\$7 500), University of Toronto	Sept. 2021
B.C. Achievement Scholarship (\$1 250), Government of British Columbia	Aug. 2021
District/Authority Scholarship (\$1 250), Government of British Columbia	Aug. 2021

STUDENT LEADERSHIP

Director of Internal Relations ; Computer Science Student Union, University of Toronto	Apr. 2023 – Apr. 2024
<ul style="list-style-type: none">Organized orientation for the ~500 undergraduate students entering the computer science streamPlanned 20+ events in collaboration with various partners in industry (such as AMD and Google) or student organizations (such as UTMIST 🔗 and WiCS 🔗)Hosted 5+ talks with professors in the Department of Computer Science at the University of Toronto	
First-Year Academic Officer ; Math Union, University of Toronto	Sept. 2021 – Apr. 2022
<ul style="list-style-type: none">Facilitated discussions between 20 mentor-mentee pairs in the <i>First-Year Mentorship Program</i> by providing guidance to the upper-year mentorsOrganized “Coffee and Chat” events which allowed for informal discussions between students and math professors	
Registered Study Group Leader ; Sidney Smith Commons, University of Toronto	Sept. 2021 – April 2022
<ul style="list-style-type: none">Led study groups for <i>Foundations of Computer Science I</i> and <i>Enriched Introduction to the Theory of Computation</i>Headed weekly meetings for first-year students that reviewed content covered in the previous week's lectureDeveloped example problems to clarify and reinforce important concepts through group discussion	

PROJECTS

Image Domain Adaption 🔗	Sept. 2023 – Dec. 2023
<ul style="list-style-type: none">Used Python Optimal Transport to compute various functions which transform the <i>EMNIST</i> dataset of handwritten digits such that its distribution and priors match those of the <i>MNIST</i> datasetFound that the accuracy of a fully-connected feed-forward classifier trained on the <i>MNIST</i> dataset was improved from 17% to 73% on the <i>EMNIST</i> dataset	
Student Response Classifier	Mar. 2023 – Apr. 2023
<ul style="list-style-type: none">Developed a 3-parameter logistic item response theory classifier in PyTorch, using alternating gradient descent for training, to predict the correctness of student answers to multiple-choice questionsObtained an accuracy of 72% on the <i>NeurIPS 2020 Education Challenge</i> dataset (within 5% of the best solution)	
Image Classifier 🔗	Dec. 2022 – Jan. 2023
<ul style="list-style-type: none">Implemented a softmax classifier with stochastic gradient descent (SGD) from scratch in C++ using only the linear algebra library Eigen3Achieved 92% accuracy on the <i>MNIST</i> dataset of handwritten digits (within 2% of the top classifier using SGD)Built in the ability to save trained weights, perform batch training, and track the training error in real-time	
Image Restoration with Convolutional Neural Networks 🔗	Sept. 2020 – June 2021
<ul style="list-style-type: none">Combined the models RIDNet and DeepDeblur using PyTorch to determine the ability of convolutional neural networks to deblur and denoise imagesArtificially generated a dataset of 5 000 noisy, blurred images using a Poisson-Gaussian noise modelDiscovered that integrating the two models offers marginal improvements over their individual performance	

TECHNICAL SKILLS

Languages	Python, C++, C, Rust, Java
Frameworks	PyTorch, TensorFlow, scikit-learn, NumPy, Pandas, SciPy, Matplotlib, MuJoCo, Eigen3
Tools	Git, shell, ssh, Unix, CMake, Anaconda, Google Colab, QEMU, Jupyter, WSL, Slurm

SELECTED COURSEWORK

Code	Title	Term
CSC2221*	Introduction to the Theory of Distributed Computing	Fall 2024
CSC324	Principles of Programming Languages	Winter 2024
CSC412†	Probabilistic Learning and Reasoning	Winter 2024
CSC413†	Neural Networks and Deep Learning	Winter 2024
CSC473	Advanced Algorithm Design	Winter 2024
MAT357	Real Analysis I	Winter 2024
APM462	Nonlinear Optimization	Fall 2023
CSC369	Operating Systems	Fall 2023
CSC420	Introduction to Image Understanding	Fall 2023
MAT354	Complex Analysis I	Fall 2023
MAT377	Mathematical Probability	Fall 2023
MAT327	Introduction to Topology	Summer 2023
CSC373	Algorithm Design, Analysis and Complexity	Winter 2023
CSC384	Introduction to Artificial Intelligence	Winter 2023
CSC438	Computability and Logic	Winter 2023
CSC463	Computational Complexity and Computability	Fall 2022
MAT344	Introduction to Combinatorics	Summer 2022

*Graduate course

†Cross-listed graduate course